
Dynamic Disequilibrium: An Agent-Based Mechanism for Bubbles without Biased Agents

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We present the *Estimated Dynamic Equilibrium Model* (EDEM), an agent-based
2 framework in which supply and demand are realisations of a coupled stochastic
3 process driven by heterogeneous, error-prone agent valuations. The framework’s
4 central technical contribution is a generative mechanism for persistent disequilib-
5 rium: when market-clearing prices are repeatedly selected from the upper tail of
6 noisy bid distributions and fed back into future valuations, expected prices drift
7 upward even with strictly zero-mean estimation errors. We derive the bias in closed
8 form for a clean special case ($\sigma(n-1)/(n+1)$ for n iid uniform bids, which is the
9 classical winner’s curse) and show via simulation that compounding the bias across
10 epochs produces exponential price growth without any behavioural assumption
11 about investor optimism. EDEM extends Miller [1977]’s divergence-of-opinion
12 theory to a fully dynamic setting and recovers Walrasian equilibrium and Miller’s
13 static premium as limiting cases. Eight controlled experiments in the Python ABM
14 framework Mesa, on a 32×32 real-estate neighbourhood, plus a 30-cell sensitivity
15 grid, demonstrate six qualitatively distinct regimes—all reachable from the same
16 agent ruleset. The result has direct implications for machine-learning valuation
17 algorithms (Zestimate-style automated valuation models): trained on historical
18 clearing prices, such systems are by construction fitting an upper-order statistic
19 of plausible sale prices rather than fair value, and at deployment time function as
20 positively-biased point estimators.

21 1 Introduction

22 The Efficient Market Hypothesis remains the dominant theory of price formation. It predicts when
23 prices reflect available information but does not provide a generative microstructural account of when
24 and why they fail to: bubbles, crashes, persistent cross-sectional return anomalies, and the stylised
25 observation that real markets very rarely sit at the textbook supply–demand fixed point. Decades
26 of empirical work have produced rival theories whose claims often contradict one another, because
27 abundant and noisy financial data can be mined to support nearly any hypothesis [Moffitt, 2017].

28 The central technical claim of this paper is that two structural features—max-bid clearing and per-
29 epoch feedback of clearing prices into agent valuations—generate persistent positive drift in realised
30 market prices even when individual agents’ estimation errors are exactly zero-mean. The mechanism
31 is order-statistic bias. Each seller selects the maximum of multiple noisy bids, and the maximum of n
32 zero-mean draws sits, in expectation, strictly above their mean; the realised maximum then feeds
33 back into the next round’s anchor value, and the bias compounds multiplicatively. No behavioural
34 assumption about investor optimism, risk aversion, or attention is required.

Thaler [1988] called this phenomenon the winner’s curse and framed it as a behavioural failure: rational bid shading should, in principle, correct the expected payment. Our point is that the curse is an order-statistic effect before it is a behavioural one, and that iterating it—feeding each round’s clearing price into the next round’s anchor—turns a one-period bias into unbounded drift in market valuations.

The implication for machine-learning valuation matters for the NeurIPS audience. An ML valuation algorithm trained on historical clearing prices—automated valuation models (AVMs), iBuying pricing models, comparable-sale algorithms—is not fitting fair value. By construction it is fitting the upper-order statistic of plausible sale prices, because the clearing price is the maximum of the bids received. At deployment time it functions as a positively-biased point estimator that may itself provide the market with a coordinating signal pulling realised winning bids further upward.

Contributions.

1. An order-statistic mechanism for bubbles without biased agents (Section 4): in closed form, for a clean special case, the expected winning bid exceeds the home value by $\sigma(n-1)/(n+1)$ when n bidders draw zero-mean errors with dispersion σ . Compounded across epochs this produces exponential price drift.
2. A unified framework (Section 3) that nests Walrasian equilibrium and Miller’s static premium as limiting cases, and admits a family of out-of-equilibrium regimes (business cycles, persistent overshoots and undershoots, runaway bubbles, constant transition) parameterised by estimation noise, balancer strength, and time-on-market.
3. Eight controlled experiments plus a 30-cell sensitivity grid (Section 6) demonstrating that the order-statistic drift produces price drift above textbook equilibrium in 100% of parameter combinations tested.
4. A discussion of implications for ML valuation systems (Section 7) connecting the model to selection-aware estimation procedures that AVM deployers should adopt.

2 Background and Related Work

Miller [1977] proposed an elegant static argument: of K investors with heterogeneous beliefs, the N most optimistic end up holding the N available shares, so the clearing price reflects the marginal optimist rather than the mean valuation. Greater divergence of opinion thus produces a price *premium*, and (because future returns are earned against the elevated price) lower subsequent returns. The half-century of empirical literature that grew up around this claim has not converged. One camp argues divergence of opinion proxies risk and predicts *higher* future returns [Williams, 1977, Mayshar, 1983, Varian, 1985, Merton, 1987, Epstein and Wang, 1994, Doukas et al., 2006]; another, anchored by Diether et al. [2002], finds the opposite empirically—lower future returns for high-dispersion stocks, consistent with Miller’s premium. Moffitt [2017] catches the spirit of the impasse: no single theory of price formation has satisfied both market practitioners and theorists. The empirical record is messy, and we will argue this is exactly what one should expect.

The static framing is the source of the empirical confusion. Miller’s model specifies a one-shot relationship between beliefs and the market-clearing price. Real markets are continuous processes in which today’s price feeds into tomorrow’s beliefs. We show that the same upper-tail-selection mechanism that produces Miller’s premium (and, separately, the winner’s curse [Thaler, 1988] in common-value auctions) generates a *dynamic* drift when iterated—and that the same underlying mechanism can produce either a premium or a discount in finite samples, which explains why empirical work has found both.

Closely related work in agent-based economics has used ABMs to study housing markets [Baptista et al., 2016, Geanakoplos et al., 2012] and behavioural finance generally [Shiller, 2003]. The contribution here is the structural identification of the selection mechanism, formalised analytically and validated empirically across a parameter sweep.

3 The EDEM Framework

A market is staged on a finite torus G of homes, with $Q_s(t)$ sellers $\mathcal{S}_t$ (one per home, immobile) and $Q_d(t)$ buyers $\mathcal{B}_t$ (mobile, walking the torus). Each home h has a fixed unobservable fair value $v^*(h)$

86 and a market value $v_t(h)$ that evolves. Each agent i valuing home h at tick t draws an estimation

$$E_{i,t}(h) = v_t(h) (1 + \varepsilon_{i,t,h}), \quad \varepsilon_{i,t,h} \stackrel{\text{iid}}{\sim} F_i(\sigma_i), \quad (1)$$

87 with $\mathbb{E}[\varepsilon] = 0$. The three-way index is essential: each *bid* is an independent draw. Throughout, σ_i is
88 expressed as a decimal (0.05 means $\pm 5\%$ maximum error); parameter tables report percentages.

89 A buyer that lands on a seller's home posts a bid $\beta_{b,s}^{(t)} = E_{b,t}(h(s))$. Sellers accumulate bids; on
90 patience timeout, the seller offers to sign with the buyer holding the best bid. The buyer commits iff

$$\beta \geq \frac{1}{|\mathcal{B}_b|} \sum_{\beta' \in \mathcal{B}_b} \beta', \quad (2)$$

91 where $\mathcal{B}_b$ is the buyer's outstanding bid set (including β). Eq. (2) is order-statistic selection on the
92 buyer's side: committing only when the offered bid is at or above the buyer's running benchmark
93 restricts purchases to the upper tail of the buyer's own bid distribution.

94 The realised market price is the output of a clearing functional $p_t = A(X_t; \theta)$ over the full market
95 state X_t . We study two instantiations. In the discrete-event *DE* variant, A is the rolling average of
96 the last W completed sales:

$$p_t = \frac{1}{m_t} \sum_{j=k(t)-m_t+1}^{k(t)} P_j^{\text{sale}}, \quad m_t = \min(W, k(t)), \quad (3)$$

97 where $k(t)$ counts completed sales up to tick t . In the speculative *EDEM* variant there are no realised
98 sales; instead each home's value updates multiplicatively at the end of every epoch of T ticks:

$$v_{t+T}(h) = v_t(h) \cdot \bar{r}_t, \quad \bar{r}_t = \frac{1}{|\mathcal{Y}_t|} \sum_{b \in \mathcal{Y}_t} \min_{s \in \mathcal{W}_b(t)} \frac{\beta_{b,s}^{(t)}}{v_t(h(s))} \quad (4)$$

99 when the set $\mathcal{Y}_t$ of buyers holding at least one winning bid is non-empty (else $\bar{r}_t = 1$). $\mathcal{W}_b(t)$ is the
100 set of sellers on which buyer b holds the winning bid.

101 The supply and demand quantities are themselves integer-valued random walks driven by the price
102 increment $\Delta p_{t-1} = p_{t-1} - p_{t-2}$:

$$Q_s(t) = Q_s(t-1) + Z_t^s, \quad \mathbb{E}[Z_t^s] = C_b \text{sgn}(\Delta p_{t-1}). \quad (5)$$

103 With $C_b > 0$ rising prices add sellers (mean-reverting); with $C_b < 0$ rising prices remove sellers
104 (trend-following); with $C_b = 0$ the population is frozen. A symmetric equation governs Q_d . A
105 population floor of one agent per side truncates the random walk in extreme regimes.

106 Setting $\sigma_i \equiv 0$ collapses Eq. (1) to universal agreement; the order-statistic drift vanishes and EDEM
107 approaches the Walrasian fixed-point equilibrium. Retaining $\sigma_i > 0$ but freezing the dynamics
108 ($T \rightarrow \infty$, single period, no value update) recovers a Miller-style static premium mechanism.

109 4 The Order-Statistic Mechanism

110 A concrete number first. With just two bidders ($n = 2$) drawing zero- mean uniform errors of width
111 $\sigma = 0.15$, the expected winning bid sits 5% above fair value—per home, per epoch. Iterate that for
112 150 epochs and the implied compound is enormous. The closed form is what the rest of this section
113 derives.

114 **Closed-form bias.** Consider a seller s receiving $n \geq 2$ bids $\beta_i = v_t(h(s))(1 + \varepsilon_i)$ with $\varepsilon_i \sim$
115 $\mathcal{U}[-\sigma, +\sigma]$ iid. The expected winning bid is

$$\mathbb{E} \left[\frac{\max_i \beta_i}{v_t(h(s))} \right] = 1 + \mathbb{E} \left[\max_i \varepsilon_i \right] = 1 + \sigma \frac{n-1}{n+1}, \quad (6)$$

116 using the order-statistic identity for the maximum of n iid uniform variates. The right-hand side
117 strictly exceeds 1 for all $n \geq 2$ and $\sigma > 0$; the gap—whether one calls it excess return or mispricing—
118 grows in both n and σ .

Relation to the winner’s curse. Eq. (6) recovers the classical winner’s curse as a static special case. In a common-value auction, the winning bidder is selected precisely because her signal is the most favourable draw among the bidder pool. Thaler [1988] treats this as a behavioural failure; rational bid shading should correct the expected payment in any one round. Whether shading fully corrects, or only partially, is downstream of the structural fact: the winning signal comes from the upper tail, full stop. EDEM’s contribution is to compound that fact—one round’s clearing price becomes the next round’s anchor—and to show that iteration alone turns a one-period bias into unbounded drift.

What Eq. (6) does and does not prove. Eq. (6) establishes that the seller-side maximum has expected ratio strictly above 1. The per-epoch update Eq. (4) compounds a related but distinct quantity $\bar{r}_t$ (the buyer-side mean of $\min_s$ winning ratios), for which we have no closed form. Two further care points apply when moving from the theorem to sample-path behaviour: (i) for an iid multiplicative process $v_T = v_0 \prod_t r_t$, the expectation $\mathbb{E}[v_T] = v_0 (\mathbb{E}[r])^T$ grows multiplicatively whenever $\mathbb{E}[r] > 1$, but the median is governed by $\mathbb{E}[\log r]$; by Jensen’s inequality $\mathbb{E}[r] > 1$ does not on its own imply $\mathbb{E}[\log r] > 0$. (ii) The fallback $\bar{r}_t = 1$ for empty-winner epochs puts a point mass at unity that the clean theorem does not capture.

Empirical bridge. Table 1 reports the distribution of $\bar{r}_t$ extracted directly from our EDEM-regime simulations, with $\bar{r}_t = v_{t+T}(h)/v_t(h)$ at every epoch boundary across all seeds. The growth condition $\mathbb{E}[\log \bar{r}_t] > 0$ holds in every regime, certifying sample-path exponential growth and closing the gap between Eq. (6) and the simulated buyer-min update.

Table 1: Empirical distribution of $\bar{r}_t$ across 1500 epochs per regime (10 seeds, 150 epochs/seed). The growth condition $\mathbb{E}[\log \bar{r}_t] > 0$ holds in every regime examined.

Regime	$\mathbb{E}[\bar{r}_t]$	$\Pr(\bar{r}_t > 1)$	$\mathbb{E}[\log \bar{r}_t]$	median $\bar{r}_t$
Run 6 ($C_b = 0, \bar{\sigma} = 0.15$)	1.0429	83.5%	+0.0411	1.044
Run 7 ($C_b = +1, \bar{\sigma} = 0.15$)	1.0104	19.7%	+0.0098	1.000
Run 7 ($C_b = -1, \bar{\sigma} = 0.15$)	1.0153	25.4%	+0.0144	1.000
Run 8 ($C_b = -1, \bar{\sigma}$ growing)	1.0287	24.9%	+0.0236	1.000

5 Implementation

EDEM is implemented in the Python ABM framework Mesa [Masad and Kazil, 2015] on a 32×32 toroidal MultiGrid, matching the wrapped patch world of the original NetLogo prototypes [Wilensky, 1999] from which this work descends. Per-cell state ($v_t(h)$, $v^*(h)$, last sale price, last sale tick) is a lightweight dataclass attached to each grid cell; buyers and sellers are `mesa.Agent` subclasses. Bids are mirrored dictionaries on each side. The rolling 25-sale market price of Eq. (3) is a `deque(maxlen = 25)`. The per-epoch update in Eq. (4) fires on a cycle-counter trick one tick before buyers reset, preserving their winning-bid flags. The balancer Eq. (5) fires $\lfloor |C_b| \rfloor$ deterministic swaps per epoch plus one Bernoulli swap with probability $|C_b| - \lfloor |C_b| \rfloor$; sign of C_b controls direction. A population floor of one agent per side prevents the balancer from killing the last agent.

The Mesa core lives in six modules under `edem/` totalling 973 lines including docstrings; the 42 pytest unit tests sit in five files of similar density. Each experiment script seeds the model deterministically and runs 10 seeds. Total reproduction time on a 2024-era laptop is approximately 25 minutes (CPU only). Code and data are released under MIT licence at a URL provided in the camera-ready version of this paper.

6 Experiments and Results

Eight controlled experiments grouped into two regimes (DE: Runs 1–5, discrete-event sales; speculative-market EDEM: Runs 6–8, multiplicative value updates) plus a 30-cell sensitivity grid (Run 9). Each experiment runs 10 seeds; figures plot the median across seeds with the 10th–90th percentile band. Full parameter tables are in Appendix B; per-run figures for runs not shown inline here are in Appendix C.

159 **The six regimes.** Three parameter axes ($\bar{\sigma}$, seller patience, C_b) and the same agent ruleset generate
160 six qualitatively distinct steady states: band-stable, business-cycle, persistent overshoot, persistent
161 undershoot, runaway bubble, and constant transition. Table 2 consolidates the end-of-run statistics;
162 we flag the one entry below the table that drives the rest of this section.

Table 2: Eight EDEM runs: parameters and observed outcomes (median across ≥ 8 seeds). DE runs report median deviation of the rolling market price from the textbook equilibrium; EDEM runs report v_t/v^* .

Run	Manipulated parameter	Outcome metric	Regime
1	$\bar{\sigma} = 0.05$, patience ≤ 50	−2% to +11% band	band-stable
2	$\bar{\sigma} = 0.25$	−17% to +34% periodic	business-cycle
3	patience ≤ 100	+14% persistent	overshoot
4	demand intercept −50	−18% persistent	undershoot
5	shock at $t = 3000$, then re-shocks	moving-target chase	transitional
6	EDEM, $C_b = 0$, $\bar{\sigma} = 0.15$	$v_t/v^* = 421\times$	runaway bubble
7	EDEM, $C_b \in \{-1, 0, +1\}$	7.7, 421, 4.2 $\times$ resp.	balancer-shaped
8	EDEM, $C_b = -1$, $\bar{\sigma}$ growing	$v_t/v^* = 30\times$	super-linear

163 **Bubble regime under no balancer.** With balancer disabled ($C_b = 0$) and $\bar{\sigma} = 0.15$, the market
164 exhibits a clean log-linear bubble: a median multiplicative gain of $\sim 421\times$ over 3000 ticks, with a
165 tight band (Fig. 1). The mechanism is exactly that of Section 4—each epoch’s average winning ratio
166 $\bar{r}_t$ is > 1 in expectation, and Eq. (4) compounds it multiplicatively. $C_b = 0$ admits all of the upward
167 bid-selection bias unfiltered, which is exactly what this configuration is meant to expose.




Figure 1: Run 6 (bubble, $C_b = 0$, $\bar{\sigma} = 0.15$). Top: median v_t/v^* (log scale) with 10–90 percentile band. Bottom: agent counts (frozen at 20 each because $C_b = 0$).

168 **Balancer-coefficient sweep.** Varying $C_b \in \{+1, 0, -1\}$ at fixed $\bar{\sigma} = 0.15$ yields three qualitatively
169 distinct trajectories (Fig. 2). The mean-reverting case ($C_b = +1$) plateaus near $4\times$. The trend-
170 following case ($C_b = -1$) reaches $\sim 7.7\times$. The no-balancer case dominates both at over $100\times$
171 either, ending near $421\times$. The intuitive guess, dating to earlier informal write-ups in this lineage, is
172 that trend-following should be the most explosive regime—there is more aggregate bidding pressure
173 when buyers swarm a shrinking pool of sellers. In the Mesa simulation the opposite holds. The
174 trend-following balancer drains sellers to a single agent, and inflation needs sellers to inflate; the

no-balancer case keeps populations balanced and admits the bid-selection bias undamped. The population panel of Fig. 2 makes the mechanism visible: the $C_b = \pm 1$ trajectories collapse into a one-vs-many state within a few hundred ticks, while $C_b = 0$ holds at the 20/20 initialisation throughout. No setting of C_b in this grid restores $v_t/v^* = 1$.




Figure 2: Run 7 (balancer sweep at $\bar{\sigma} = 0.15$). The three regimes are overlaid; $C_b = 0$ produces the strongest bubble. The population panel shows the one-vs-many state the strong balancers induce.

Sensitivity to C_b and $\bar{\sigma}$. A 5×6 grid sweeps $C_b \in \{-1, -0.5, 0, +0.5, +1\}$ and $\bar{\sigma} \in \{0.05, 0.10, 0.15, 0.20, 0.25, 0.30\}$, with 5 seeds per cell at 1500 ticks (Fig. 3). Two patterns emerge. First, the gap widens monotonically with $\bar{\sigma}$ —larger divergence of opinion produces a larger multiplier, in every column. Second, every cell of the grid has $v_t/v^* > 1$, and the $C_b = 0$ column dominates each row. Of the 30 cells, 13 exceed $10\times$ within 1500 ticks. The order-statistic drift is universal across the parameter grid we explored, not a corner case.

Other regimes. DE Runs 1–5 produce the band-stable, business-cycle, overshoot, undershoot, and transitional regimes listed in Table 2. Run 4 (low agent density): mean realised price $p \approx 41$ versus textbook $p^* = 50$, an 18% shortfall, with episodes reaching $p \approx 37$ (26% at the trough). Run 5 (shocked market): the realised price never settles when demand is toggled every 2000 ticks. The figures for these runs are in Appendix C.

7 Discussion: Implications for ML Valuation

Bid-selection asymmetry, not estimation bias, drives bubbles. The estimation function Eq. (1) has zero mean error: an agent’s expected valuation is exactly the home’s current value. Yet the EDEM speculative regime produces unbounded upward drift even with this unbiased estimator. The mechanism is not behavioural bias but a statistical asymmetry: each seller selects the maximum bid received, and the maximum of N noisy estimates of a fair value is, in expectation, larger than the fair value itself. Eq. (4) compounds this asymmetry multiplicatively across epochs. This is the dynamic generalisation of the classical winner’s curse [Thaler, 1988].

Implication for ML valuation systems. The same mechanism operates in any pricing system whose inputs are anchored to recent market-clearing prices. A machine-learning valuation algorithm trained on historical sale prices is, by construction, fitting the upper-order statistic of plausible sale prices for a similar asset—the maximum of the relevant noisy estimates—not their mean. Deployed at




Figure 3: Run 9 (sensitivity heatmap). Median terminal v_t/v^* over a 5×6 ($C_b, \bar{\sigma}$) grid, log colour scale. Red squares mark cells corresponding to Runs 6 and 7. Every cell exceeds 1.0; 43% exceed 10.0.

scale, such an algorithm functions as a positively-biased point estimator that may provide the market with a coordinating signal pulling realised winning bids further upward, feeding back into the next training cycle.

The fix is structural, not a hyperparameter knob. A defensible deployment would predict the median rather than the mean of plausible sale prices, use a censored-data likelihood that accounts for the offer-acceptance threshold, or substitute an explicitly rank-aware estimator that targets fair value instead of winning-bid value. The Q3 2021 disclosures of Zillow’s algorithmic-iBuying losses [Zillow Group, Inc., 2021] document one episode in which a clearing-price-anchored valuation system at scale incurred substantial losses; the public record is consistent with the mechanism described here, though we do not advance EDEM as a complete account of that specific event.

Why empirical premium/discount studies disagree. Each empirical sample of the divergence-of-opinion premium is a finite-window observation along a coupled stochastic process governed by σ and C_b . Two studies sampling distinct positions on distinct trajectories may both correctly observe opposite signs. The unifying claim from EDEM is not that one camp is right and the other wrong, but that the static framing has been asking the wrong question; the relevant quantity is the trajectory of $\bar{r}_t$, not its sign at a snapshot.

8 Limitations and Conclusion

EDEM models a single asset class with no leverage, no macroeconomic environment, and non-adaptive (non-learning) agents. The 30-cell sensitivity grid covers a defensible range of ($C_b, \bar{\sigma}$) but does not explore extreme $|C_b|$ values where the population dynamics qualitatively change. Run 5’s shock-recovery scenario uses a 12,000-tick budget that is too short for the market to reach its post-shock floor before the patience boost fires. Bayesian-belief and reinforcement-learning extensions of the agent rule set (Appendix D) are sketched but not implemented.

The contribution we believe travels beyond the EDEM setting is the mechanism. Bubbles, persistent disequilibrium, and the empirical contradictions in the divergence-of-opinion literature do not require behavioural bias to explain. Repeated upper-tail selection, fed back into subsequent valuations, is sufficient. For the NeurIPS audience the practical consequence is concrete: an ML valuation model trained on historical clearing prices is by construction fitting an upper-order statistic, and selection-aware estimation is a deployment-time requirement, not an academic nicety.

References

- Anonymous Author(s). Disequilibria in markets with stochastic processes: An agent-based model approach. Working paper, 2018. (Anonymised for double-blind review.), 2018a.
- Anonymous Author(s). Estimated dynamic equilibrium model: Supply and demand as a sample path of a stochastic process. Conference paper, 2018. (Anonymised for double-blind review.), 2018b.
- Rafa Baptista, J. Doyne Farmer, Marc Hinterschweiger, Katie Low, Daniel Tang, and Arzu Uluc. Macroprudential policy in an agent-based model of the UK housing market. *Bank of England Staff Working Paper*, (619), 2016.
- Karl B. Diether, Christopher J. Malloy, and Anna Scherbina. Differences of opinion and the cross section of stock returns. *The Journal of Finance*, 57(5):2113–2141, 2002. doi: 10.1111/0022-1082.00490.
- John A. Doukas, Chansog (Francis) Kim, and Christos Pantzalis. Divergence of opinion and equity returns. *Journal of Financial and Quantitative Analysis*, 41(3):573–606, 2006. doi: 10.1017/S0022109000002544.
- Larry G. Epstein and Tan Wang. Intertemporal asset pricing under Knightian uncertainty. *Econometrica*, 62(2):283–322, 1994.
- John Geanakoplos, Robert Axtell, J. Doyne Farmer, Peter Howitt, Benjamin Conlee, Jonathan Goldstein, Matthew Hendrey, Nathan M. Palmer, and Chun-Yi Yang. Getting at systemic risk via an agent-based model of the housing market. *American Economic Review*, 102(3):53–58, 2012. doi: 10.1257/aer.102.3.53.
- Volker Grimm, Steven F. Railsback, Christian E. Vincenot, Uta Berger, Cara Gallagher, Donald L. DeAngelis, Bruce Edmonds, Jiaqi Ge, Jarl Giske, Jürgen Groeneveld, Alice S. A. Johnston, Alexander Milles, Jacob Nabe-Nielsen, J. Gareth Polhill, Viktoriia Radchuk, Marie-Sophie Rohwäder, Richard A. Stillman, Jan C. Thélin, and Sven Berger. The ODD protocol for describing agent-based and other simulation models: A second update to improve clarity, replication, and structural realism. *Journal of Artificial Societies and Social Simulation*, 23(2):7, 2020. doi: 10.18564/jasss.4259.
- David Masad and Jacqueline Kazil. Mesa: An agent-based modeling framework. In *Proceedings of the 14th Python in Science Conference (SciPy 2015)*, pages 51–58, 2015. doi: 10.25080/Majora-7b98e3ed-009.
- Joram Mayshar. On divergence of opinion and imperfections in capital markets. *The American Economic Review*, 73(1):114–128, 1983.
- Robert C. Merton. A simple model of capital market equilibrium with incomplete information. *The Journal of Finance*, 42(3):483–510, 1987. doi: 10.1111/j.1540-6261.1987.tb04565.x.
- Edward M. Miller. Risk, uncertainty, and divergence of opinion. *The Journal of Finance*, 32(4):1151–1168, 1977. doi: 10.1111/j.1540-6261.1977.tb03317.x.
- Steven D. Moffitt. *The Strategic Analysis of Financial Markets, Volume I: Framework*. World Scientific, 2017. Earlier circulated as “Why Markets are Inefficient: A Gambling Theory of Financial Markets For Practitioners and Theorists,” Feb. 22, 2017.
- Robert J. Shiller. From efficient markets theory to behavioral finance. *Journal of Economic Perspectives*, 17(1):83–104, 2003. doi: 10.1257/089533003321164967.
- Richard H. Thaler. Anomalies: The winner’s curse. *Journal of Economic Perspectives*, 2(1):191–202, 1988. doi: 10.1257/jep.2.1.191.
- Hal R. Varian. Divergence of opinion in complete markets: A note. *The Journal of Finance*, 40(1):309–317, 1985. doi: 10.1111/j.1540-6261.1985.tb04951.x.
- Uri Wilensky. NetLogo. <http://ccl.northwestern.edu/netlogo/>, 1999. Center for Connected Learning and Computer-Based Modeling, Northwestern University, Evanston, IL.

- 277 Joseph T. Williams. Capital asset prices with heterogeneous beliefs. *Journal of Financial Economics*,
278 5(2):219–239, 1977.
- 279 Zillow Group, Inc. Quarterly report (form 10-q) for the quarter ended september 30, 2021: Wind down
280 of Zillow Offers. U.S. Securities and Exchange Commission, EDGAR filing accession 0001617640-
281 21-000087, November 2021. URL [https://www.sec.gov/Archives/edgar/data/1617640/](https://www.sec.gov/Archives/edgar/data/1617640/000161764021000087/z-20210930.htm)
282 [000161764021000087/z-20210930.htm](https://www.sec.gov/Archives/edgar/data/1617640/000161764021000087/z-20210930.htm).

A ODD Protocol

This appendix documents EDEM in the standard ODD protocol for agent-based models [Grimm et al., 2020].

Purpose and patterns. EDEM characterises the conditions under which an agent-based real-estate market reaches a stable equilibrium price, and identifies the mechanisms that prevent equilibrium when they do not. Patterns the model reproduces: stable band-bounded equilibria; endogenous business cycles; persistent shifts in the realised price relative to the textbook equilibrium under altered patience or density; multiplicative price drift in the absence of a balancer.

Entities, state variables, scales. Entities: Buyers and Sellers (`mesa.Agent` subclasses), Homes (per-cell dataclass on a 32×32 toroidal grid). Buyer state: epsilon (estimation-error bound), delay, current bid, heading; EDEM variant adds the lowest-bid-to-value ratio and a yellow (winning) flag. Seller state: epsilon, patience timer, current ask price, dictionary of received bids; EDEM adds number of bids, total bids, best-bid-to-value, best-bid-to-true-value. Home state: market price, last-sale tick, last-sale price, fair value v^* , current value v . Model state: rolling window of last 25 sale prices (DE), epoch counter (EDEM), cross-population epsilon (EDEM). One spatial unit = one home; one tick is the smallest temporal unit. DE balance period is 100 ticks; EDEM epoch is 20.

Process overview. Per tick: (i) each agent steps once, in randomised order over both classes; a Buyer’s step posts bids, then wiggles, then steps forward; a Seller’s step processes patience and best-bid logic; (ii) the model invokes the balancer (DE: every 100 ticks; EDEM: every 20 ticks via the cycle-counter trick that preserves buyers’ winning-bid flags); (iii) the data collector samples model-level reporters.

Design concepts. *Basic principles:* EDEM extends the divergence-of-opinion premise of Miller [1977] to a multi-period setting; agents act on noisy estimates of an unobservable fair value. *Emergence:* the price level, agent population dynamics, and macroscopic regimes emerge from local interaction; none of these phenomena are coded in directly. *Adaptation:* sellers lower ask prices when patience elapses without high-enough bids (the only adaptive behaviour). Buyers do not adapt. *Objectives:* sellers maximise realised sale price subject to patience. Buyers commit only to offers whose realised bid is at or above their own running benchmark (Eq. (2))—a selection rule, not a price-minimisation objective. *Learning, prediction:* none in the baseline model. *Sensing:* buyers sense whether a seller is on their patch; sellers sense the bids in their own bid table; neither senses the global market price directly. *Interaction:* direct via mirrored-dictionary updates; indirect via the average ask price and rolling sale-price average influencing newly-spawned agents. *Stochasticity:* agent placement, buyer heading, epsilon draws, per-bid error draws, victim selection, Bernoulli draws for fractional C_b swaps—all from a single seeded `numpy.random.Generator`. *Observation:* per-tick model reporters stacked into Parquet datasets, one per seed.

Initialisation. DE: spawn `equi_qnty` sellers at random unique cells with patience $\sim \mathcal{U}[0, \text{init_patience})$ and ask price drawn from $\mathcal{U}[-\bar{\sigma}, +\bar{\sigma}]$ around p^* ; spawn `equi_qnty` buyers at random cells. EDEM: 20 sellers and 20 buyers, all homes initialised with $v_0(h) = v^*(h) = 100$.

Input data and submodels. No external input data; EDEM is closed. Estimation function: Eq. (1). Bid acceptance: Eq. (2). Market price: Eq. (3) (DE) and Eq. (4) (EDEM). Balancer: Eq. (5) in finite-population form with population floor of one.

B Parameter Tables

Common parameters across all runs. World size 32×32 toroidal; seeds per run ≥ 8 (typically 10); bid-acceptance rule `netlogo` (the rule actually coded in the original 2018 NetLogo source, matching Eq. (2)).

DE runs 1–5. Supply $\hat{Q}_s(p) = a_s + 0.5p$, $a_s = 0$. Demand $\hat{Q}_d(p) = a_d - 0.5p$, $a_d = 100$ (Runs 1–3, 5) or 50 (Run 4) or $\{125, 75\}$ alternating (Run 5 Scenario B). Maximum valuation error $\bar{\sigma}$: 0.05

(Runs 1, 3–5) or 0.25 (Run 2). Maximum patience: 50 (Runs 1, 2, 4) or 100 (Run 3) or 50 $\rightarrow$ 165 at $t = 7000$ (Run 5 Scenario A). Balance period $T_B = 100$. Sale window $W = 25$. Ticks per seed: 20,000 (Runs 1–4), 12,000 (Run 5).

EDEM runs 6–8. Initial $\bar{\sigma}$: 0.15 (Runs 6, 7) or 0.05 (Run 8). $\bar{\sigma}$ growth per epoch: 0 (Runs 6, 7) or +0.005 (Run 8). Balancer C_b : 0 (Run 6); $\{+1, 0, -1\}$ (Run 7); -1 (Run 8). Initial buyers / sellers: 20/20. Epoch length 20 ticks. Ticks per seed: 3000. True value $v^*(h) = 100$ uniform across all homes.

Run 9 sensitivity grid. $C_b \in \{-1, -0.5, 0, +0.5, +1\}$; $\bar{\sigma} \in \{0.05, 0.10, 0.15, 0.20, 0.25, 0.30\}$. 5 seeds per cell, 1500 ticks each.

C All Experiment Figures




Figure 4: Run 1 – Stable equilibrium with low valuation error and moderate patience.

D Extensions and Future Work

The framework of Section 3 is deliberately minimal so that extensions can be plugged in without rewriting the agent skeleton. Five directions are especially fruitful and we sketch each below.

Bias-corrected clearing. The bid-selection asymmetry embedded in the max-bid clearing rule cannot be cancelled by any linear balancer within the explored grid. Two replacements for the clearing functional A in Eq. (3) would address it directly: (i) a winsorised mean replacing $\bar{r}_t$ in Eq. (4); (ii) a selection-aware estimator substituting the median or lower-percentile of all bids (not only winning bids) for $\bar{r}_t$. Either is a one-line edit to the reference implementation.

Realtors and commission structures. The DE seller is a strict utility-maximiser whose only friction is patience. Real-world residential transactions are intermediated by realtors whose payoff is a fraction of the sale price. A realtor extension would replace the seller agent with an aggregated listing agent whose patience and ask-price strategies optimise expected commission across a portfolio. This is a natural setting for empirical calibration: realtor commission rates and listing durations are publicly observable, unlike the patience parameter of the bare DE seller.




Figure 5: Run 2 – High valuation error ($\bar{\sigma} = 0.25$) generates endogenous business cycles.




Figure 6: Run 3 – Doubling seller patience to 100 pushes the market price about 14% above the textbook equilibrium and sustains it there.

354 **Open vs. blind auctions.** EDEM as posed is a blind-bid model. Open-bid auctions substantially
 355 change the bid distribution by exposing buyers to one another's valuations and would likely amplify
 356 the bid-selection asymmetry.

357 **Construction and the supply side.** A construction-market variant inverts EDEM's agent roles:
 358 contractors bid downward to win projects, with the lowest-bid contractor winning. The selection
 359 asymmetry then operates in reverse, biasing realised contract prices below the population fair value.




Figure 7: Run 4 – Halving the demand intercept reduces equilibrium quantity to 25; the realised market price settles about 18% below the textbook equilibrium.

Adaptive agents. EDEM agents are non-learning. A Bayesian extension where each agent maintains a posterior over fair value and updates after every observed sale would converge, in the limit, toward bias-corrected estimation. A reinforcement-learning variant with a buy-low/sell-high reward would discover and exploit the bid-selection asymmetry; this would be a computational analogue to the Zillow-Offers iBuying problem [Zillow Group, Inc., 2021] and would offer a stress test of any proposed clearing-rule fix.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction state the order-statistic mechanism, the closed-form bound $\sigma(n-1)/(n+1)$, the experimental scope (eight runs plus a 30-cell sensitivity grid), the six identified regimes, and the implication for ML valuation systems. Each claim is substantiated in the body.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?




Figure 8: Run 5 – Two shock scenarios. Top: a single demand shock at $t = 3000$ followed by a patience boost at $t = 7000$. Bottom: a transitional market in which demand toggles every 2000 ticks.

Answer: [Yes]

Justification: Section 8 explicitly states the limitations: single asset class, no leverage, no macroeconomic environment, non-adaptive (non-learning) agents, finite parameter grid that does not cover extreme $|C_b|$, and a truncated shock-recovery horizon in Run 5.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an impor-




Figure 9: Run 8 – Trend-following balancer ($C_b = -1$) combined with linearly-growing $\bar{\sigma}$ ($0.05 \rightarrow 0.80$). The trajectory is sustained super-linear rather than visibly double-exponential.

tant role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: The single theoretical result (Eq. 6) is derived under explicit iid uniform-error assumptions stated in Section 4. The paper distinguishes what the closed form proves (seller-side maximum) from what the simulation certifies (the buyer-side $\min_s$ update); the empirical bridge in Table 1 closes the remaining gap with $\mathbb{E}[\log \bar{r}_t] > 0$ in every regime.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.

- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 5 describes the Mesa implementation. Appendix B lists every parameter for every run. The bid-acceptance and value-update rules are stated formally as Eqs. 2–4. Deterministic seeding, per-seed Parquet outputs, and the build script are documented in the released repository.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Code is released under MIT licence and the per-tick model reporters for every run are released as Parquet datasets. The repository URL is provided in the camera-ready version of this paper (omitted here to preserve double-blind anonymity).

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.

- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: All parameters appear in Appendix B: world size, supply/demand intercepts and slopes, $\bar{\sigma}$ range, patience bounds, balance period, sale window, epoch length, and per-run seed count. There is no train/test split or optimiser since the work is simulation-based, not learning-based.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: All trajectory figures report the median across 10 seeds (5 seeds for Run 9) with the 10th–90th percentile band as variability; this is stated in Section 6 and in figure captions. Per-cell statistics in Table 1 aggregate over 1,500 epoch boundaries per regime; the reported $\mathbb{E}[\log \bar{r}_t] > 0$ condition is the relevant sample-path-growth statistic for a multiplicative process.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).

- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: All experiments are CPU-only and reproduce in approximately 25 minutes total on a 2024-era laptop; Run 9 (the longest single script) takes about 4 minutes. The implementation requires no GPU, no cluster, and approximately 1 GB of memory. This is stated in Section 5.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research involves no human subjects, no datasets containing personally identifying information, no pre-trained models, and no scraped data. The paper conforms with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 7 discusses the broader implication that ML valuation algorithms trained on historical clearing prices inherit upper-order-statistic bias; the negative societal impact (potentially amplifying market-wide mispricing in automated valuation models deployed at scale) is made explicit, and selection-aware mitigations are sketched (predicting median rather than mean; censored-data likelihoods; rank-aware estimators).

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases an agent-based-model implementation and the resulting simulation datasets, neither of which is a high-misuse-risk artifact (no pre-trained language model, no image generator, no scraped data). No safeguards are required.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Mesa [Masad and Kazil, 2015] (Apache-2 licence) is the only third-party code asset and is cited. The Grimm et al. ODD protocol is cited as the published source for Appendix A. The 2018 NetLogo prototypes that seeded the design are cited (Anonymous Author(s) [2018a,b]) and preserved verbatim in the released codebase as historical record.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.

- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: The released Mesa codebase, the 9 figure PDFs, and the per-seed Parquet datasets are documented in the repository with a top-level README, per-experiment scripts, and 42 pytest unit tests. Code is MIT-licensed; data and paper are CC-BY-4.0.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: The paper involves no crowdsourcing or research with human subjects.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: The paper involves no human subjects research.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: Large language models were used to assist with drafting, editing, mathematical verification, and code review during paper preparation; they were not part of the core methodology. All claims, equations, references, and experimental code were independently verified by the authors. A pre-submission audit was performed specifically to check for hallucinated references, meta-comments, and placeholders; the artifacts of that audit are in the released repository.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.